

Rough Draft

Tyler Bontrager, Ganesh Singh

2022-11-02

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.8      v dplyr    1.0.10
## v tidyverse 1.2.1     v stringr  1.4.1
## v readr   2.1.2      vforcats  0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()

# IMPORTING DATASETS
tuition_cost <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/datasets/2022-11-02/tuition_cost.csv')

## # Rows: 2973 Columns: 10
## -- Column specification -----
## Delimiter: ","
## chr (5): name, state, state_code, type, degree_length
## dbl (5): room_and_board, in_state_tuition, in_state_total, out_of_state_tuit...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
tc = tuition_cost

tuition_income <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/datasets/2022-11-02/tuition_income.csv')

## # Rows: 209012 Columns: 7
## -- Column specification -----
## Delimiter: ","
## chr (4): name, state, campus, income_lvl
## dbl (3): total_price, year, net_cost
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
ti = tuition_income

salary_potential <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/datasets/2022-11-02/salary_potential.csv')

## # Rows: 935 Columns: 7
## -- Column specification -----
## Delimiter: ","
## chr (2): name, state_name
## dbl (5): rank, early_career_pay, mid_career_pay, make_world_better_percent, ...
```

```

## 
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
sp = salary_potential

historical_tuition <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/main/2023-03-28/historical_tuition.csv')

## # A tibble: 270 × 4
##   type      year tuition_type tuition_cost
##   <chr>     <dbl>    <chr>        <dbl>
## 1 Private  1960      In-State       10350
## 2 Private  1960      Out-of-State  11660
## 3 Private  1960      In-State       11318
## 4 Private  1960      Out-of-State  13680
## 5 Private  1960      In-State       10350
## 6 Private  1960      Out-of-State  11660
## 7 Private  1960      In-State       11318
## 8 Private  1960      Out-of-State  13680
## 9 Private  1960      In-State       10350
## 10 Private 1960      Out-of-State  11660
## # ... with 2,693 more rows, 1 more variable: tuition_type <chr>, and
## #   abbreviated factor levels: In-State, Out-of-State
## #   and 1 more variable: tuition_cost <dbl>

## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
ht = historical_tuition

## # A tibble: 50655 × 5
##   name      state category total_enrollment enrollment
##   <chr>     <chr>    <chr>        <dbl>        <dbl>
## 1 Aaniih ~ Mont~ MT      Public~ 2 Year          10350
## 2 Abilene ~ Texas TX     Private~ 4 Year         11660
## 3 Abraham ~ Geor~ GA     Public~ 2 Year          10350
## 4 Academy ~ Minn~ MN     For ~ 2 Year          11318
## 5 Academy ~ Cali~ CA     For ~ 4 Year          11318
## 6 Adams St~ Colo~ CO    Public~ 4 Year          10350
## 7 Adelphi ~ New ~ NY    Private~ 4 Year         11660
## 8 Adironda~ New ~ NY    Public~ 2 Year          11318
## 9 Adrian C~ Mich~ MI    Private~ 4 Year         11318
## 10 Advanced~ Virg~ VA   For ~ 2 Year          11318
## # ... with 2,963 more rows, 1 more variable: category <chr>, and
## #   abbreviated factor levels: Public~, For ~, Private~, and
## #   1 more variable: enrollment <dbl>

## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
ds = diversity_school

## # A tibble: 2,973 × 11
##   name      state state-1 type degree~2 room_~3 in_st~4 in_st~5 out_o~6 out_o~7
##   <chr>     <chr>    <chr>    <chr>    <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Aaniih ~ Mont~ MT      Publ~ 2 Year          NA        2380      2380      2380      2380
## 2 Abilene ~ Texas TX     Priv~ 4 Year         10350     34850     45200     34850     45200
## 3 Abraham ~ Geor~ GA     Publ~ 2 Year          8474      4128      12602     12550     21024
## 4 Academy ~ Minn~ MN     For ~ 2 Year          NA        17661     17661     17661     17661
## 5 Academy ~ Cali~ CA     For ~ 4 Year         16648      27810     44458     27810     44458
## 6 Adams St~ Colo~ CO    Publ~ 4 Year          8782      9440      18222     20456     29238
## 7 Adelphi ~ New ~ NY    Priv~ 4 Year         16030     38660     54690     38660     54690
## 8 Adironda~ New ~ NY    Publ~ 2 Year          11660      5375      17035     9935      21595
## 9 Adrian C~ Mich~ MI    Priv~ 4 Year         11318     37087     48405     37087     48405
## 10 Advanced~ Virg~ VA   For ~ 2 Year          NA        13680     13680     13680     13680
## # ... with 2,963 more rows, 1 more variable: degFactor <fct>, and
## #   abbreviated factor levels: state_code, degree_length, room_and_board,
## #   in_state_tuition, in_state_total, out_of_state_tuition,
## #   out_of_state_total
## #   and 1 more variable: totalCost <dbl>

gatheredtc = tcFactored %>%
  filter(type != 'Other') %>%
  gather(key = "in_out", value = "totalCost", c(in_state_total, out_of_state_total))

```

```
gatheredtc
```

```
## # A tibble: 5,944 x 11
##   name      state state_code type degree_length room_and_board in_st_tuition in_st_total out_of_state_tuition out_of_state_total degFactor
##   <chr>     <chr>    <chr> <chr>        <dbl>           <dbl>          <dbl>       <dbl>           <dbl>          <dbl>
## 1 Aaniih N~ Mont~ MT      Publ~ 2 Year        NA        2380        2380 2 Year      in_st~
## 2 Abilene C~ Texas TX     Priv~ 4 Year      10350      34850      34850 4 Year      in_st~
## 3 Abraham B~ Geor~ GA     Publ~ 2 Year        8474       4128       12550 2 Year      in_st~
## 4 Academy C~ Minn~ MN     For ~ 2 Year        NA       17661       17661 2 Year      in_st~
## 5 Academy o~ Cali~ CA     For ~ 4 Year      16648       27810      27810 4 Year      in_st~
## 6 Adams Sta~ Colo~ CO     Publ~ 4 Year        8782       9440       20456 4 Year      in_st~
## 7 Adelphi U~ New ~ NY    Priv~ 4 Year      16030      38660      38660 4 Year      in_st~
## 8 Adirondac~ New ~ NY    Publ~ 2 Year      11660       5375       9935 2 Year      in_st~
## 9 Adrian Co~ Mich~ MI    Priv~ 4 Year      11318      37087      37087 4 Year      in_st~
## 10 Advanced ~ Virg~ VA   For ~ 2 Year        NA       13680       13680 2 Year      in_st~
## # ... with 5,934 more rows, 1 more variable: totalCost <dbl>, and abbreviated
## #   variable names 1: state_code, 2: degree_length, 3: room_and_board,
## #   4: in_state_tuition, 5: out_of_state_tuition, 6: degFactor
str(tcFactored)

## tibble [2,973 x 11] (S3:tbl_df/tbl/data.frame)
## $ name : chr [1:2973] "Aaniih Nakoda College" "Abilene Christian University" "Abrah...
## $ state : chr [1:2973] "Montana" "Texas" "Georgia" "Minnesota" ...
## $ state_code : chr [1:2973] "MT" "TX" "GA" "MN" ...
## $ type : chr [1:2973] "Public" "Private" "Public" "For Profit" ...
## $ degree_length : chr [1:2973] "2 Year" "4 Year" "2 Year" "2 Year" ...
## $ room_and_board : num [1:2973] NA 10350 8474 NA 16648 ...
## $ in_state_tuition : num [1:2973] 2380 34850 4128 17661 27810 ...
## $ in_state_total : num [1:2973] 2380 45200 12602 17661 44458 ...
## $ out_of_state_tuition: num [1:2973] 2380 34850 12550 17661 27810 ...
## $ out_of_state_total : num [1:2973] 2380 45200 21024 17661 44458 ...
## $ degFactor : Factor w/ 3 levels "2 Year","4 Year",...: 1 2 1 1 2 2 2 1 2 1 ...
head(tcFactored)

## # A tibble: 6 x 11
##   name      state state_code type degree_length room_and_board in_st_tuition in_st_total out_of_state_tuition out_of_state_total degFactor
##   <chr>     <chr>    <chr> <chr>        <dbl>           <dbl>          <dbl>       <dbl>           <dbl>          <dbl>
## 1 Aaniih N~ Mont~ MT      Publ~ 2 Year        NA        2380        2380 2380        2380
## 2 Abilene C~ Texas TX     Priv~ 4 Year      10350      34850      45200      34850 45200
## 3 Abraham B~ Geor~ GA     Publ~ 2 Year        8474       4128       12602       12550 21024
## 4 Academy C~ Minn~ MN     For ~ 2 Year        NA       17661       17661 17661       17661
## 5 Academy o~ Cali~ CA     For ~ 4 Year      16648       27810      44458       27810 44458
## 6 Adams Sta~ Colo~ CO     Publ~ 4 Year        8782       9440       18222       20456 29238
## # ... with 1 more variable: degFactor <fct>, and abbreviated variable names
## #   1: state_code, 2: degree_length, 3: room_and_board, 4: in_state_tuition,
## #   5: in_state_total, 6: out_of_state_tuition, 7: out_of_state_total
# Time to explore the data!
table(gatheredtc$state_code,gatheredtc$degFactor)

##
##      2 Year 4 Year Other
## AK      2     10     0
## AL     42     66     0
```

##	AR	48	44	0
##	AS	2	0	0
##	AZ	46	22	0
##	CA	238	270	0
##	CO	36	40	0
##	CT	28	44	0
##	DC	0	16	0
##	DE	8	10	0
##	FL	66	110	0
##	GA	58	100	0
##	GU	2	0	0
##	HI	16	12	0
##	IA	36	68	0
##	ID	8	18	0
##	IL	104	146	0
##	IN	36	88	0
##	KS	50	54	0
##	KY	30	58	0
##	LA	16	52	0
##	MA	42	144	0
##	MD	32	58	0
##	ME	18	36	0
##	MI	60	96	0
##	MN	66	76	0
##	MO	46	100	0
##	MS	30	34	0
##	MT	22	22	0
##	NC	118	116	0
##	ND	18	18	0
##	NE	20	46	0
##	NH	14	28	0
##	NJ	42	66	0
##	NM	28	20	0
##	NV	8	12	0
##	NY	116	326	0
##	OH	94	160	0
##	OK	30	50	0
##	OR	30	50	0
##	PA	62	258	0
##	PR	12	70	0
##	RI	2	20	0
##	SC	46	68	0
##	SD	10	26	0
##	TN	34	90	0
##	TX	134	164	0
##	UT	8	20	0
##	VA	60	98	0
##	VI	0	2	0
##	VT	6	32	0
##	WA	66	54	0
##	WI	62	72	0
##	WV	18	42	0
##	WY	14	2	0

```

bystate = gathereddtc %>%
  group_by(state) %>%
  mutate(freq = n()) %>%
  summarize(numSchools = sum(freq)) %>%
  mutate(prop= numSchools/sum(numSchools)) %>%
  arrange(desc(prop))
bystate

```

```

## # A tibble: 51 x 3
##   state      numSchools     prop
##   <chr>        <int>    <dbl>
## 1 California    258064 0.211
## 2 New York     195364 0.159
## 3 Pennsylvania  102400 0.0835
## 4 Texas        88804  0.0724
## 5 Ohio          64516  0.0526
## 6 Illinois      62500  0.0510
## 7 North Carolina 54756  0.0447
## 8 Massachusetts 34596  0.0282
## 9 Florida       30976  0.0253
## 10 Georgia      24964  0.0204
## # ... with 41 more rows

```

This table shows that California has a very high number of schools proportional to other states in the country.

```
prop.table(table(gathereddtc$degFactor))
```

```

##
##   2 Year    4 Year    Other
## 0.3768506 0.6231494 0.0000000
table(gathereddtc$state)

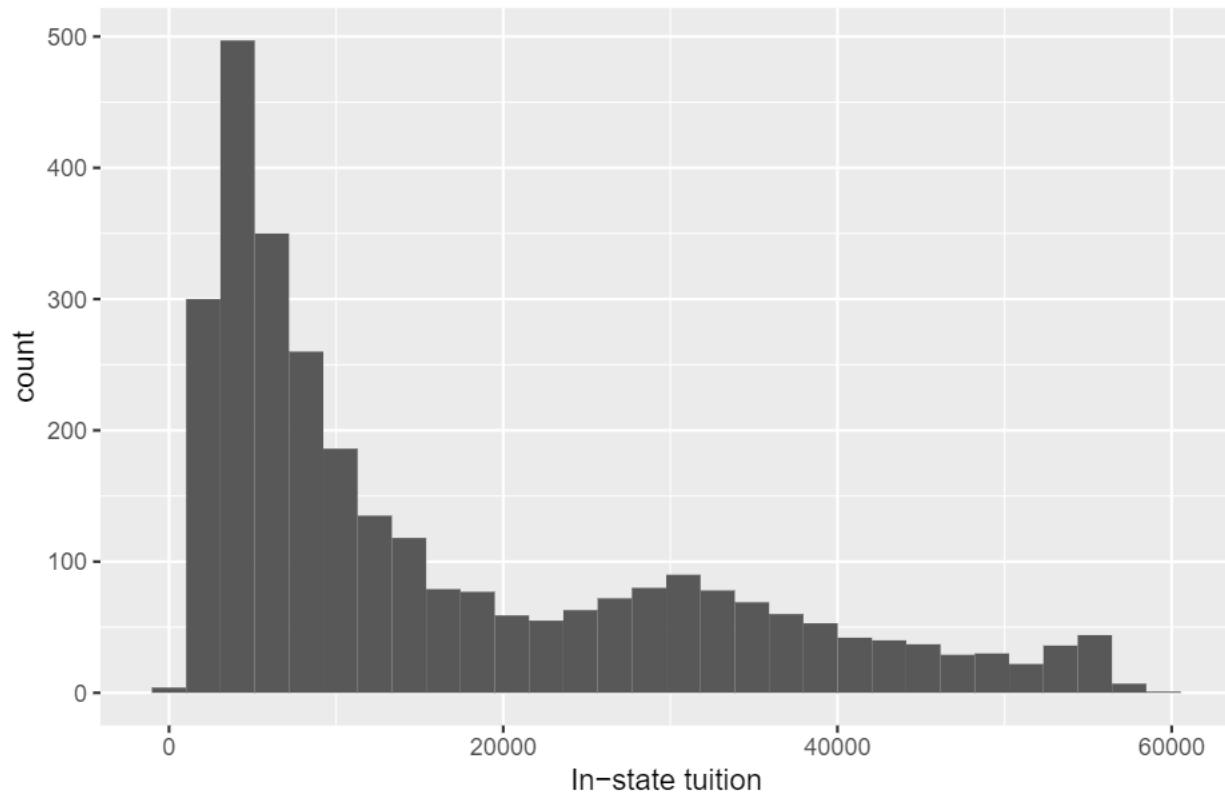
```

	Alabama	Alaska	Arizona	Arkansas	California
Colorado	108	12	68	92	508
Hawaii	76	72	18	176	158
Kansas	28	26	250	124	104
Massachusetts	104	88	68	54	90
Montana	186	156	142	64	146
New Mexico	44	Nebraska	Nevada	New Hampshire	New Jersey
Oklahoma	48	66	20	42	108
South Dakota	80	442	234	North Dakota	Ohio
Tennessee	36	80	320	36	254
Virginia	158	Texas	Pennsylvania	Rhode Island	South Carolina
Washington	120	124	298	22	114
West Virginia	120	120	60	Utah	Vermont
Wisconsin	120	120	60	28	38
Wyoming	120	120	60	Wisconsin	Wyoming
	158	120	60	134	16

```
ggplot(tcFactored, aes(x=in_state_tuition)) + geom_histogram() +
  ggtitle("Distribution of tuition charged by schools in the U.S.")+
  xlab("In-state tuition")
```

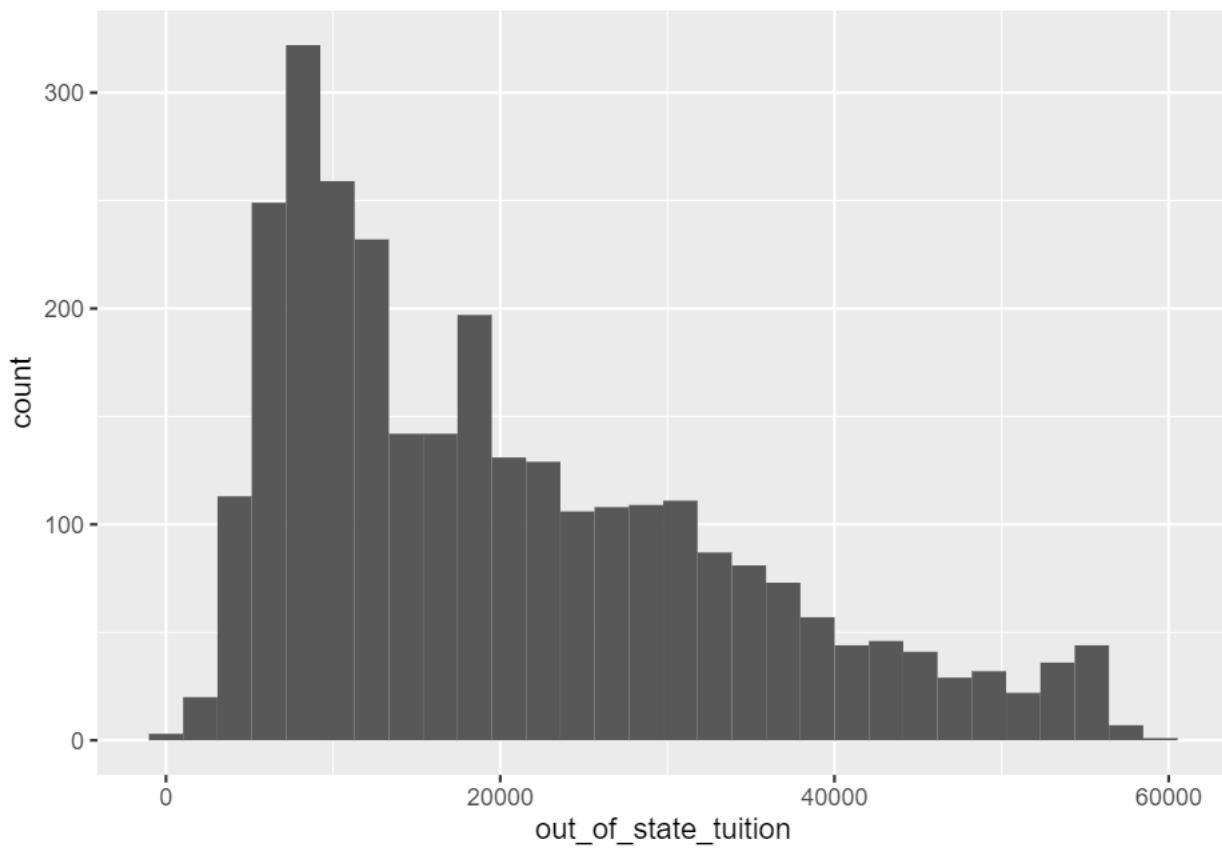
```
## `stat_bin()` using `bins = 30` . Pick better value with `binwidth`.
```

Distribution of tuition charged by schools in the U.S.



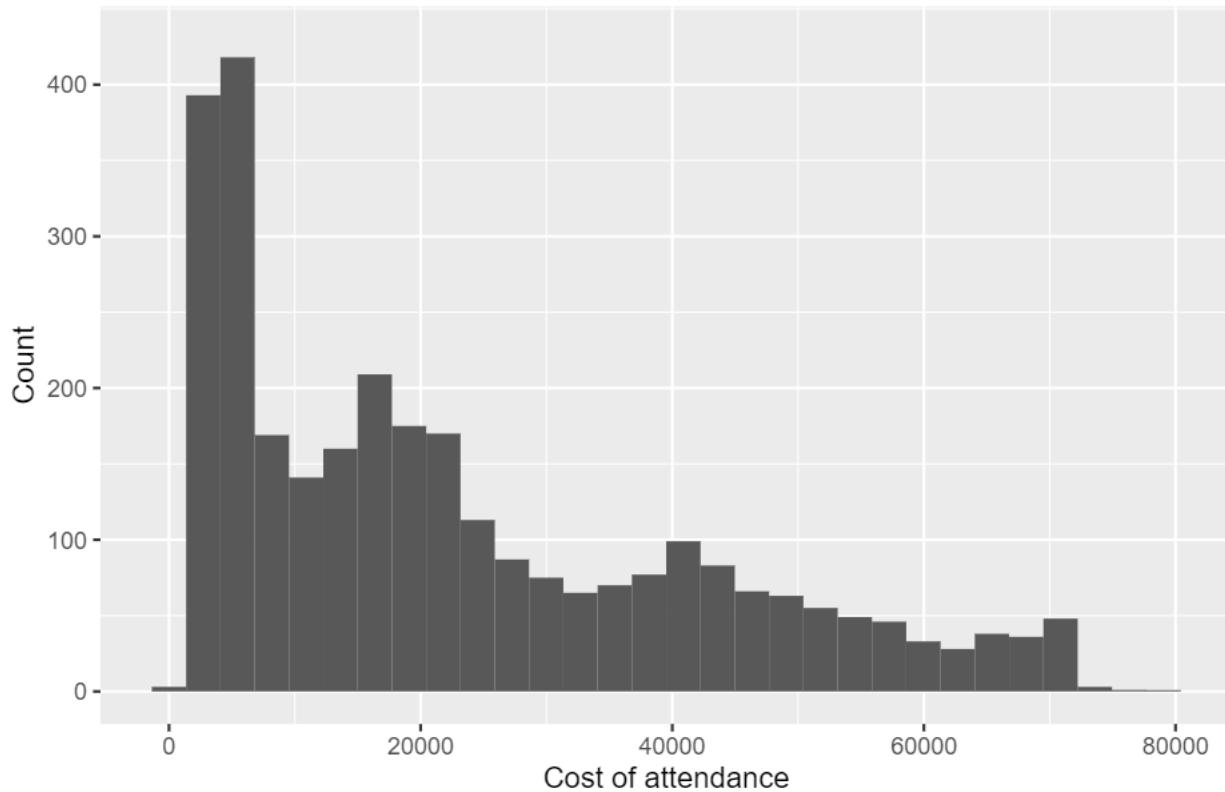
```
ggplot(tcFactored, aes(x=out_of_state_tuition))+geom_histogram()
```

```
## `stat_bin()` using `bins = 30` . Pick better value with `binwidth`.
```



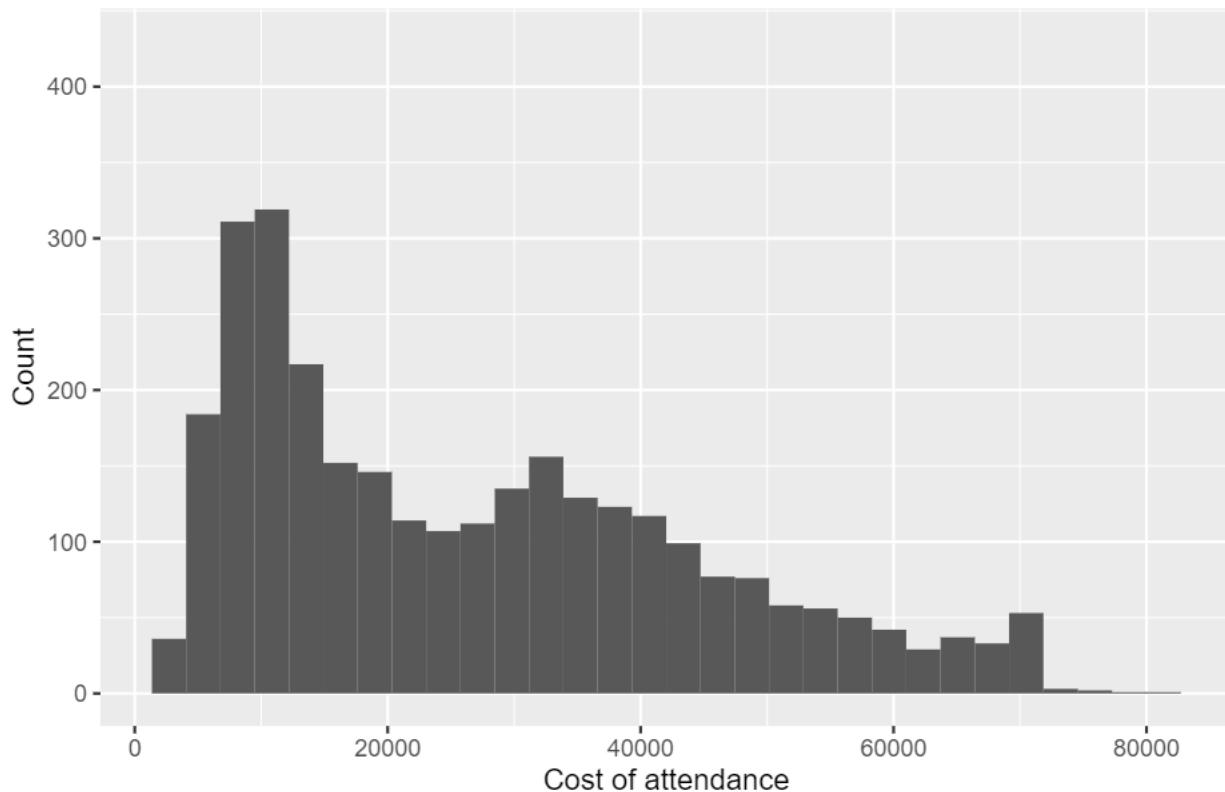
```
ggplot(tcFactored, aes(x=in_state_total))+geom_histogram()+expand_limits(x=80000,y=430) +  
  ggtitle("Total cost of attendance for in-state students")+ # for the main title  
  xlab("Cost of attendance")+ # for the x axis label  
  ylab("Count") # for the y axis label  
  
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Total cost of attendance for in-state students



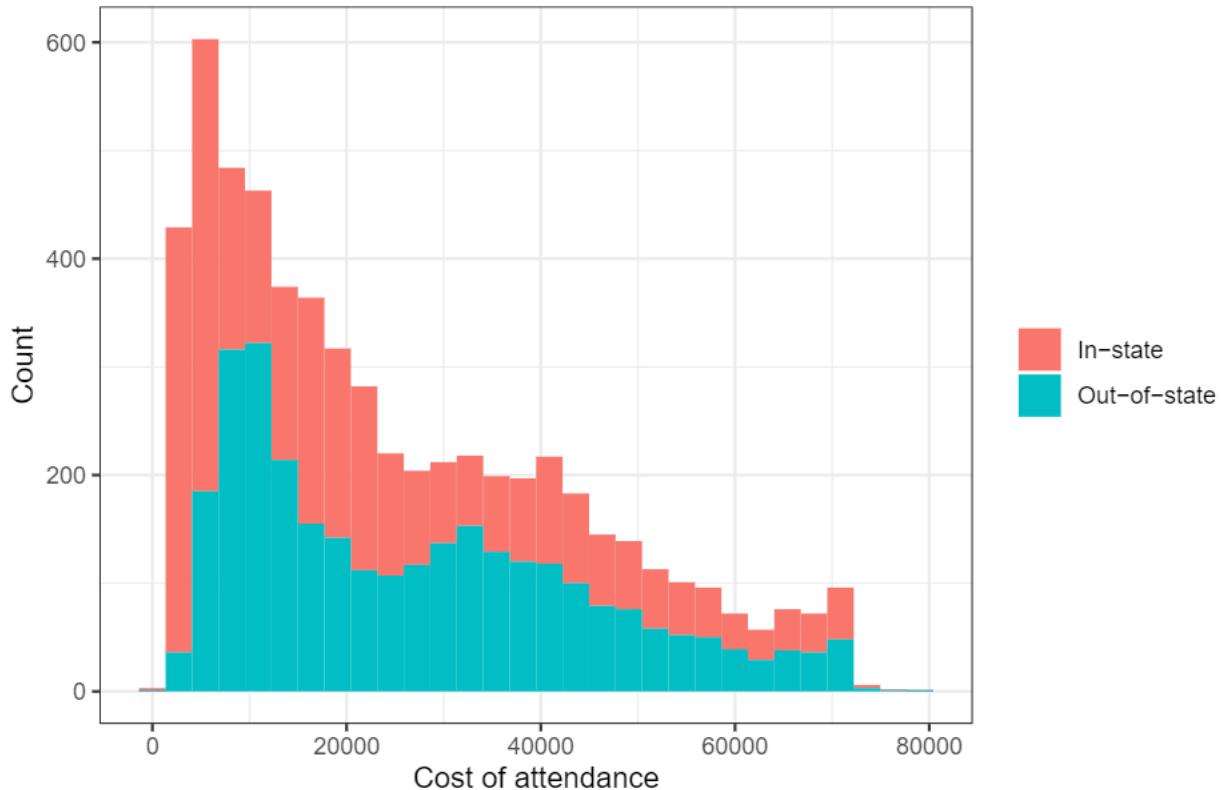
```
ggplot(tcFactored, aes(x=out_of_state_total))+geom_histogram()+expand_limits(x=80000,y=430) +  
  ggtitle("Total cost of attendance for out-of-state students")+ # for the main title  
  xlab("Cost of attendance")+ # for the x axis label  
  ylab("Count") # for the y axis label  
  
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Total cost of attendance for out-of-state students



```
ggplot(gatheredtc, aes(x=totalCost, fill=in_out))+geom_histogram()+expand_limits(x=80000, y=430) +  
  ggtitle("Total cost of attendance for students by residence") + # for the main title  
  xlab("Cost of attendance") + # for the x axis label  
  ylab("Count") + # for the y axis label  
  theme_bw() + theme(  
    legend.title = element_blank(),  
  ) + scale_fill_discrete(name = "Student Residence", labels = c("In-state", "Out-of-state"))  
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Total cost of attendance for students by residence

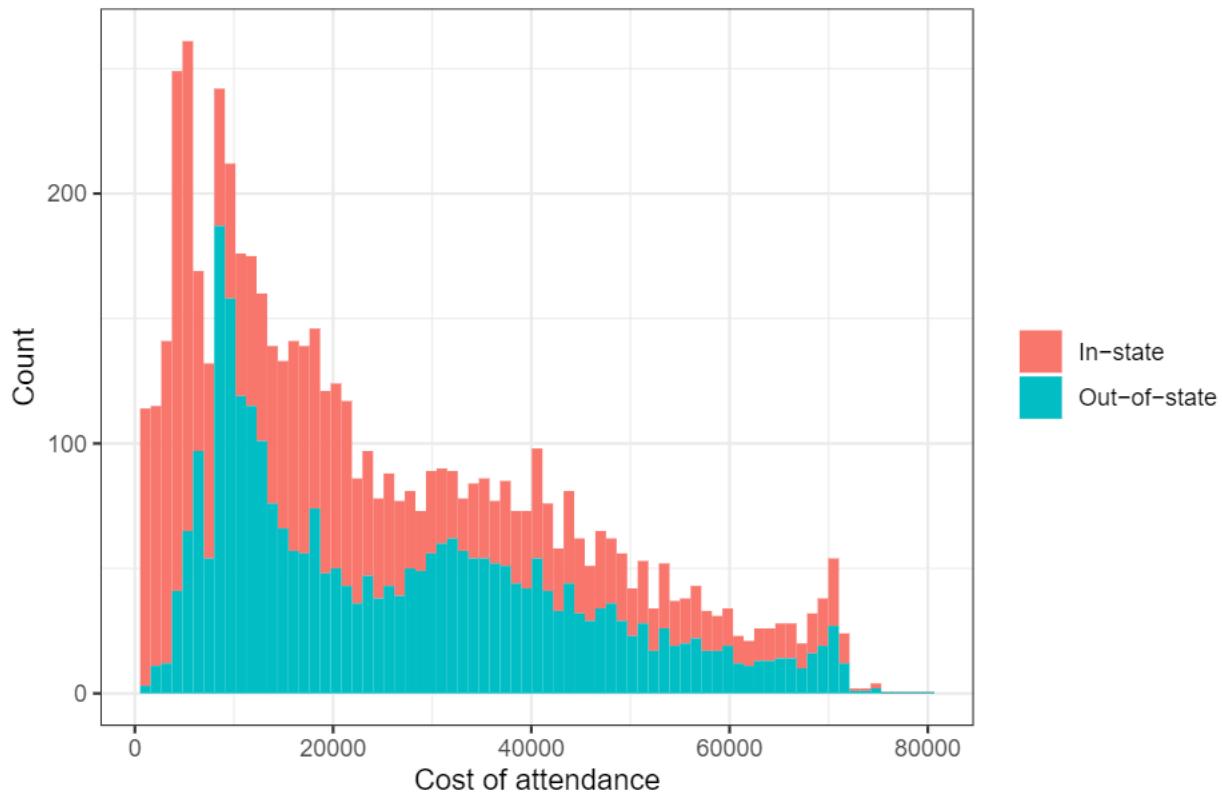


```
#ggtitle(label) # for the main title  
#xlab(label) # for the x axis label  
#ylab(label) # for the y axis label  
#labs(...) # for the main title, axis labels and legend titles
```

As the above plots show, it's clear that the distributions are skewed to the right which means that expensive schools are generally less common. It's interesting to see that both of these seem to have similar shapes, and a hint of evidence for a slight bimodal distribution.

```
ggplot(gatheredtc, aes(x=totalCost, fill=in_out))+geom_histogram(bins=75)+expand_limits(x=80000) +  
  ggtitle("Total cost of attendance for students by residence") + # for the main title  
  xlab("Cost of attendance") + # for the x axis label  
  ylab("Count") + # for the y axis label  
  theme_bw() + theme(  
    legend.title = element_blank(),  
  ) + scale_fill_discrete(name = "Student Residence", labels = c("In-state", "Out-of-state"))
```

Total cost of attendance for students by residence



Upon further inspection by increasing the bin number, the shape becomes more distinct. The second mode is mostly just a bump for the out-of-state group, but something interesting appears in the in-state group! Is there a cause for this disruption?

```
tcInStateSummr = tcFactored %>%
  group_by(degFactor) %>%
  summarize(median(in_state_total))

tcOutStateSummr = tcFactored %>%
  group_by(degFactor) %>%
  summarize(median(out_of_state_total))

tcInStateSummr

## # A tibble: 3 x 2
##   degFactor `median(in_state_total)`
##   <fct>          <dbl>
## 1 2 Year        4972.
## 2 4 Year       28287
## 3 Other         8448

tcOutStateSummr

## # A tibble: 3 x 2
##   degFactor `median(out_of_state_total)`
##   <fct>          <dbl>
## 1 2 Year      10291
## 2 4 Year     34888
## 3 Other       14640
```

This is a simple calculation of the median for 2-year and 4-year schools for total cost to out-of-state students.

```
tcFours = tcFactored %>%
  filter(degFactor=="4 Year")
tcFours

## # A tibble: 1,852 x 11
##   name      state state~1 type degree~2 room_~3 in_st~4 in_st~5 out_o~6 out_o~7
##   <chr>     <chr> <chr> <chr> <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Abilene ~ Texas TX     Priv~ 4 Year     10350    34850    45200    34850    45200
## 2 Academy ~ Cali~ CA     For ~ 4 Year    16648    27810    44458    27810    44458
## 3 Adams St~ Colo~ CO    Publ~ 4 Year     8782     9440     18222    20456    29238
## 4 Adelphi ~ New ~ NY    Priv~ 4 Year    16030    38660    54690    38660    54690
## 5 Adrian C~ Mich~ MI    Priv~ 4 Year    11318    37087    48405    37087    48405
## 6 Adventis~ Flor~ FL    Priv~ 4 Year     4200     15150    19350    15150    19350
## 7 Agnes Sc~ Geor~ GA    Priv~ 4 Year    12330    41160    53490    41160    53490
## 8 Alabama ~ Alab~ AL    Publ~ 4 Year     8379     9698     18077    17918    26297
## 9 Alabama ~ Alab~ AL    Publ~ 4 Year     5422     11068    16490    19396    24818
## 10 Alaska B~ Alas~ AK   Priv~ 4 Year    5700     9300     15000    9300     15000
## # ... with 1,842 more rows, 1 more variable: degFactor <fct>, and abbreviated
## #   variable names 1: state_code, 2: degree_length, 3: room_and_board,
## #   4: in_state_tuition, 5: in_state_total, 6: out_of_state_tuition,
## #   7: out_of_state_total

tcTwos = tcFactored %>%
  filter(degFactor=="2 Year")

tc4Y00S_Summary = tcFours%>%
  summarise(count_4Y00S=n(),
            min=min(tcFours$out_of_state_total, na.rm=TRUE),
            Q1=quantile(tcFours$out_of_state_total, prob=0.25,na.rm=TRUE),
            med=median(tcFours$out_of_state_total, na.rm=TRUE), #or quantile(AQI,prob=0.5,na.rm=TRUE)
            Q3=quantile(tcFours$out_of_state_total, prob=0.75,na.rm=TRUE),
            max=max(tcFours$out_of_state_total, na.rm=TRUE))

tc4YIS_Summary = tcFours%>%
  summarise(count_4YIS=n(),
            min=min(tcFours$in_state_total, na.rm=TRUE),
            Q1=quantile(tcFours$in_state_total, prob=0.25,na.rm=TRUE),
            med=median(tcFours$in_state_total, na.rm=TRUE), #or quantile(AQI,prob=0.5,na.rm=TRUE)
            Q3=quantile(tcFours$in_state_total, prob=0.75,na.rm=TRUE),
            max=max(tcFours$in_state_total, na.rm=TRUE))

tc2Y00S_Summary = tcTwos%>%
  summarise(count_2Y00S=n(),
            min=min(tcTwos$out_of_state_total, na.rm=TRUE),
            Q1=quantile(tcTwos$out_of_state_total, prob=0.25,na.rm=TRUE),
            med=median(tcTwos$out_of_state_total, na.rm=TRUE), #or quantile(AQI,prob=0.5,na.rm=TRUE)
            Q3=quantile(tcTwos$out_of_state_total, prob=0.75,na.rm=TRUE),
            max=max(tcTwos$out_of_state_total, na.rm=TRUE))

tc2YIS_Summary = tcTwos%>%
  summarise(count_2YIS=n(),
            min=min(tcTwos$in_state_total, na.rm=TRUE),
            Q1=quantile(tcTwos$in_state_total, prob=0.25,na.rm=TRUE),
```

```

med=median(tcTwos$in_state_total, na.rm=TRUE), #or quantile(AQI,prob=0.5,na.rm=TRUE)
Q3=quantile(tcTwos$in_state_total, prob=0.75,na.rm=TRUE),
max=max(tcTwos$in_state_total, na.rm=TRUE))

tc4Y00S_Summary

## # A tibble: 1 x 6
##   count_4Y00S   min    Q1    med    Q3   max
##       <int> <dbl> <dbl> <dbl> <dbl>
## 1      1852  1430  24951 34888 46670 75003

tc4YIS_Summary

## # A tibble: 1 x 6
##   count_4YIS   min    Q1    med    Q3   max
##       <int> <dbl> <dbl> <dbl> <dbl>
## 1      1852  1430  18199 28287 44846. 75003

tc2Y00S_Summary

## # A tibble: 1 x 6
##   count_2Y00S   min    Q1    med    Q3   max
##       <int> <dbl> <dbl> <dbl> <dbl>
## 1      1120  1376  8196. 10291 13598 68640

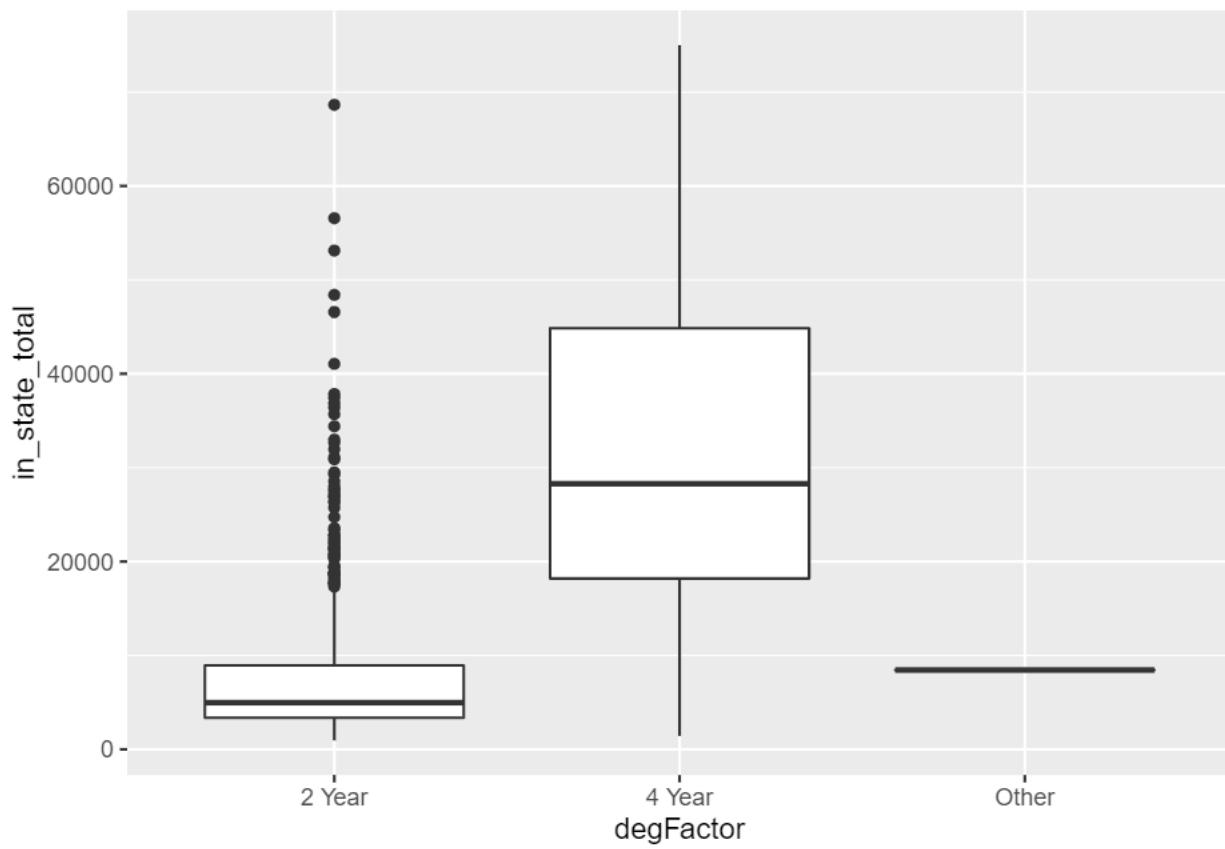
tc2YIS_Summary

## # A tibble: 1 x 6
##   count_2YIS   min    Q1    med    Q3   max
##       <int> <dbl> <dbl> <dbl> <dbl>
## 1      1120  962  3364. 4972. 8946 68640

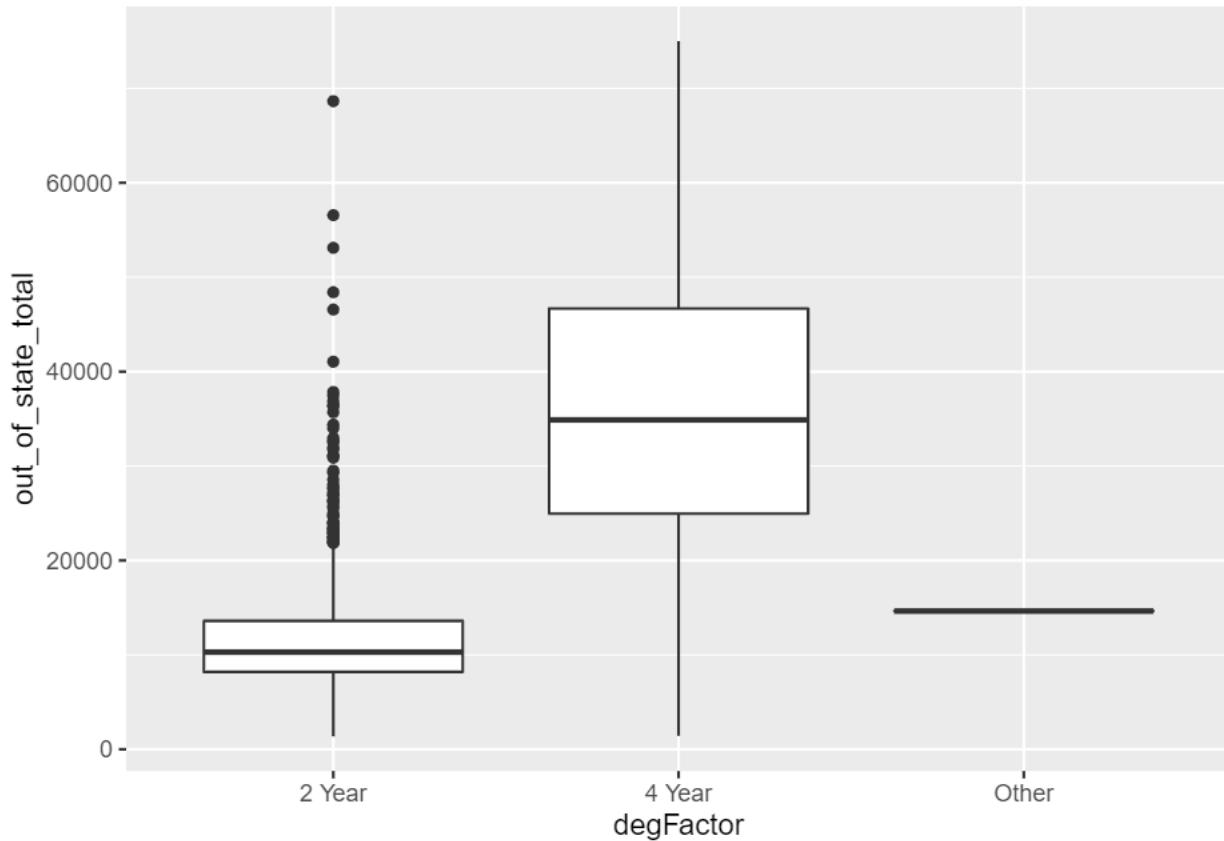
```

These are the 5-number summaries for each of the categorical variables of interest.

```
ggplot(tcFactored, aes(x = degFactor, y = in_state_total)) + # ggplot function
  geom_boxplot()
```



```
ggplot(tcFactored, aes(x = degFactor, y = out_of_state_total)) + # ggplot function  
  geom_boxplot()
```

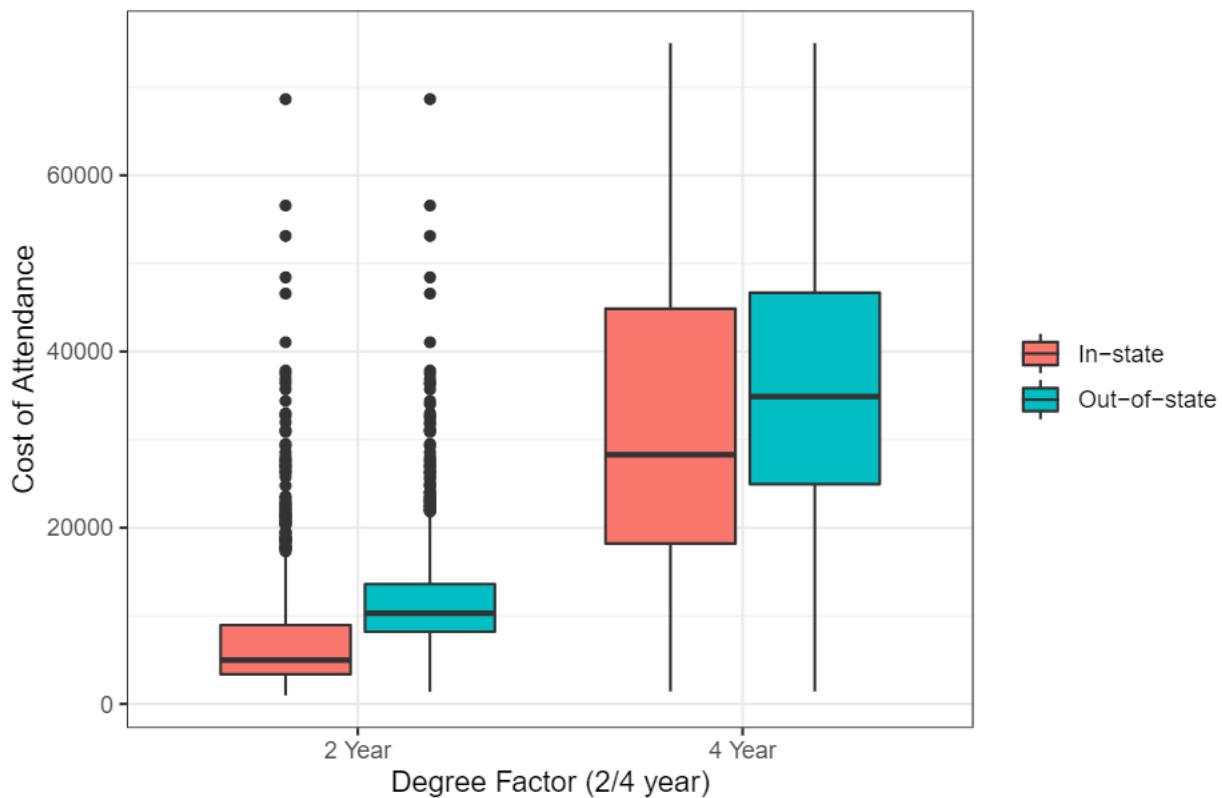


These box plots (couldn't figure out how to make an overlayed boxplot with both in/out of state variables) show a clear difference in the general cost between 2-year and 4-year institutions, and that out-of-state students generally pay more.

```
#ggplot(tcFactored, aes(x=tcInStateSummr$degFactor, fill=tcInStateSummr$in_state_total)) +
#  geom_histogram( color="#e9ecef", alpha=0.6, position = 'identity') +
#  scale_fill_manual(values=c("#69b3a2", "#404080"))

ggplot(gatheredtgc, aes(x = degFactor, y = totalCost, fill=in_out)) + # ggplot function
  geom_boxplot()+
  ggtitle("Breakdown of total cost of attendance by degree factor")+ # for the main title
  xlab("Degree Factor (2/4 year)")+ # for the x axis label
  ylab("Cost of Attendance")+ # for the y axis label
  theme_bw()+theme(
    legend.title = element_blank(),
  ) + scale_fill_discrete(name = "Student Residence", labels = c("In-state", "Out-of-state"))
```

Breakdown of total cost of attendance by degree factor




There is clearly a difference here between how much students should expect to pay given their residency status, but it isn't as absurdly significant as we were anticipating given that we hear from high school guidance counselors, specifically about 4-year institutions. Therefore, we should look for another potential explanation for the contribution to higher costs of attendance for some students.

Total Cost of Attendance by Student Residence and College Type

```
ggplot(gatheredtc,aes(x=totalCost, color=in_out)) +
  geom_histogram() +
  expand_limits(x=80000) +
  ggtitle("Distribution of cost-of-attendance by student residence")+
  xlab("Cost of attendance ($)") +
  ylab("Count") +
  theme_bw() + theme(
    legend.title = element_blank(),
  ) + scale_fill_discrete(name = "Student Residence", labels = c("In-state", "Out-of-state"))

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```


Distribution of cost-of-attendance by student residence



```
ggplot(gathereddtc,aes(x=totalCost, color=type)) +
  geom_histogram() +
  expand_limits(x=80000) +
  ggtitle("Distribution of cost-of-attendance by college type")+
  xlab("Cost of attendance ($)") +
  ylab("Count") +
  theme_bw() + theme(
    legend.title = element_blank())


## `stat_bin()` using `bins = 30` . Pick better value with `binwidth` .
```

Distribution of cost-of-attendance by college type




```
ggplot(gatheredtc,aes(x=totalCost, color=in_out)) +
  geom_density() +
  expand_limits(x=80000) +
  ggtitle("Distribution of cost-of-attendance by student residence")+
  xlab("Cost of attendance ($)") +
  ylab("Density") +
  theme_bw() + theme(
    legend.title = element_blank(),
  ) + scale_fill_discrete(name = "Student Residence", labels = c("In-state", "Out-of-state"))
```

Distribution of cost-of-attendance by student residence



```
ggplot(gatheredtc,aes(x=totalCost, color=type)) +  
  geom_density() +  
  expand_limits(x=80000) +  
  ggtitle("Distribution of cost-of-attendance by college type") +  
  xlab("Cost of attendance ($)") +  
  ylab("Density") +  
  theme_bw() + theme(  
    legend.title = element_blank())
```

Distribution of cost-of-attendance by college type



```
# ggplot(tcFactored, aes(x=in_state_total))+geom_histogram()+expand_limits(x=80000,y=430) +
#   ggtitle("Total cost of attendance for in-state students")+
#   xlab("Cost of attendance")+
#   ylab("Count")
#
# ggplot(tcFactored, aes(x=out_of_state_total))+geom_histogram()+expand_limits(x=80000,y=430) +
#   ggtitle("Total cost of attendance for out-of-state students")+
#   xlab("Cost of attendance")+
#   ylab("Count")
#
#ggttitle(label) # for the main title
#xlab(label) # for the x axis label
#ylab(label) # for the y axis label
#labs(...) # for the main title, axis labels and legend titles
```

This is a density plot that shows how high the cost of attendance is for schools across the country relative to each other given institution type.

Time to explore the data!

```
# Commenting out ggplot stuff to do dplyr first
#ggplot(ti,aes(x=year, y=total_price)) + geom_point()

#this is the median cost of attendance for instate/outstate

median_IN_COA <- tc %>%
```

```

filter(degree_length=='4 Year') %>%
group_by(state_code)%>%
summarize(median_instate_COA = median(in_state_total))

median_IN_COA

## # A tibble: 53 x 2
##   state_code median_instate_COA
##   <chr>          <dbl>
## 1 AK              17017
## 2 AL              18646
## 3 AR              19023
## 4 AZ              25037
## 5 CA              30416
## 6 CO              20976.
## 7 CT              46455
## 8 DC              50702.
## 9 DE              26542
## 10 FL             23352
## # ... with 43 more rows

median_OUT_COA <- tc %>%
filter(degree_length=='4 Year') %>%
group_by(state_code)%>%
summarize(median_outstate_COA = median(out_of_state_total))

median_OUT_COA

## # A tibble: 53 x 2
##   state_code median_outstate_COA
##   <chr>          <dbl>
## 1 AK              28604
## 2 AL              27880
## 3 AR              23709
## 4 AZ              37190
## 5 CA              36103
## 6 CO              36096
## 7 CT              48656.
## 8 DC              50702.
## 9 DE              30700
## 10 FL             32000
## # ... with 43 more rows

# mutate(mean_instate_COA=mean(in_state_tuition)) %>%
# mutate(mean_outofstate_COA=mean(out_of_state_tuition))%>

#In the following graph, I want to find out whether colleges with higher STEM enrollment tend to cost more

jointisp = ti %>%
left_join(sp) %>%
group_by(stem_percent)%>%
summarize(medianNet=median(net_cost))

## Joining, by = "name"

```

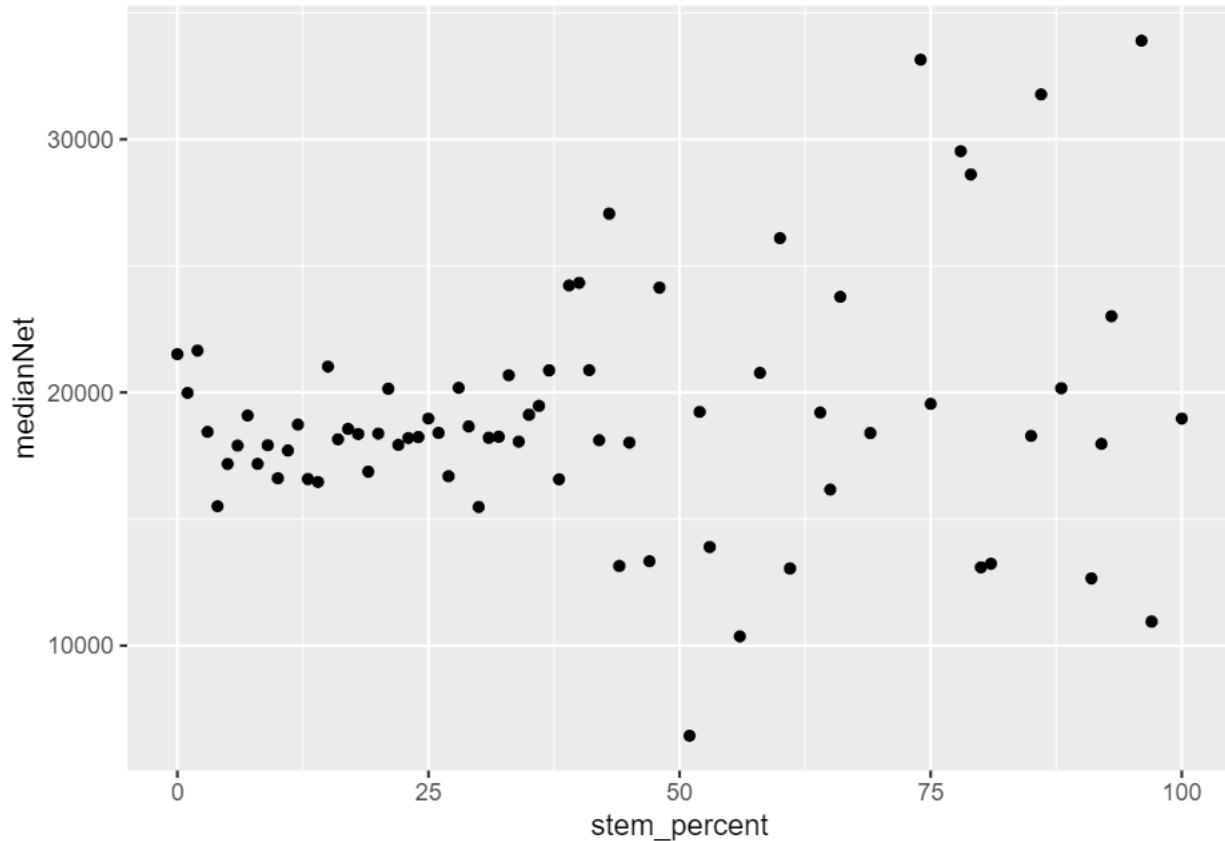
```

jointisp

## # A tibble: 75 x 2
##   stem_percent medianNet
##       <dbl>      <dbl>
## 1 0     21516.
## 2 1     19981.
## 3 2     21653
## 4 3     18443.
## 5 4     15509
## 6 5     17175
## 7 6     17907
## 8 7     19088
## 9 8     17178
## 10 9     17915
## # ... with 65 more rows
ggplot(jointisp, aes(stem_percent, medianNet)) + geom_point()

## Warning: Removed 1 rows containing missing values (geom_point).

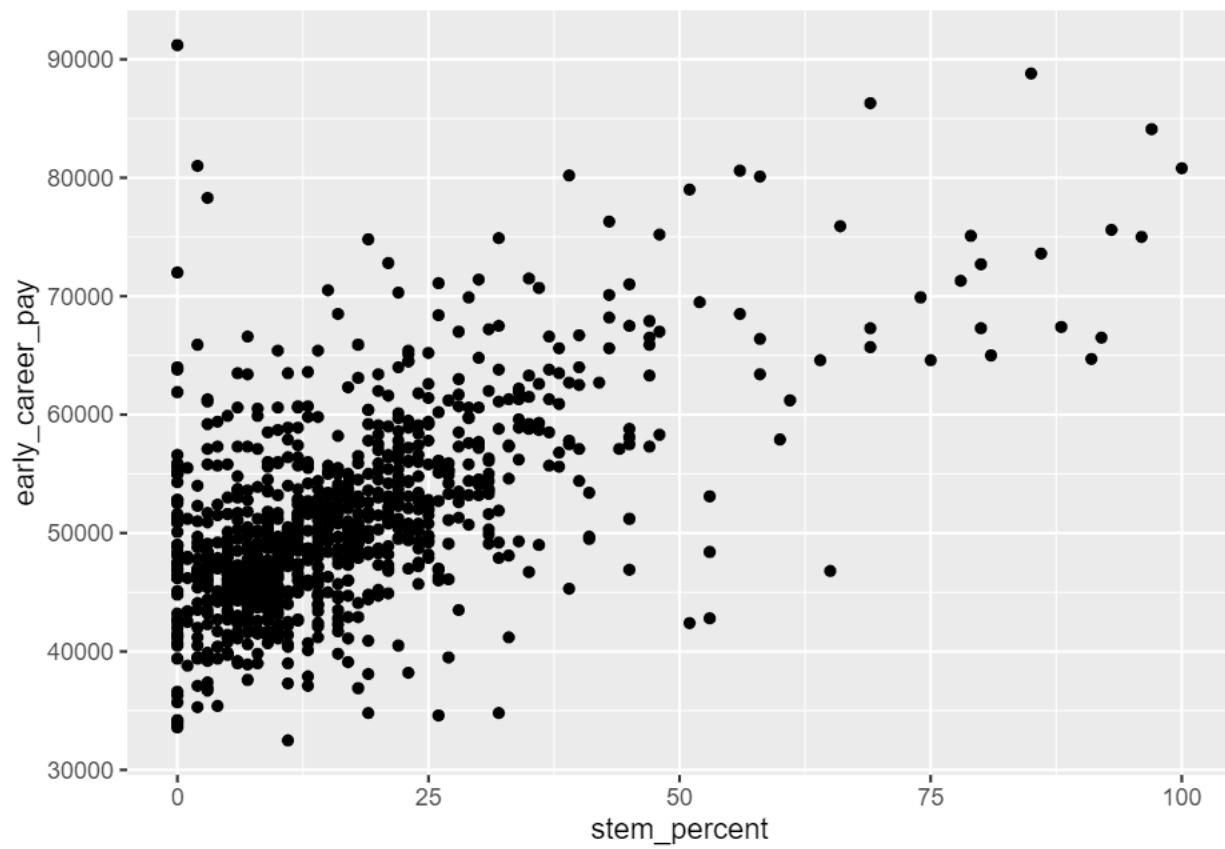
```



#My conclusion is that there is no association between higher STEM enrollment and median net cost.

In this graph I want to see whether a higher STEM enrollment has a high association with early career pay. In a later graphic, I will want to see whether the trend keeps for mid-career pay.

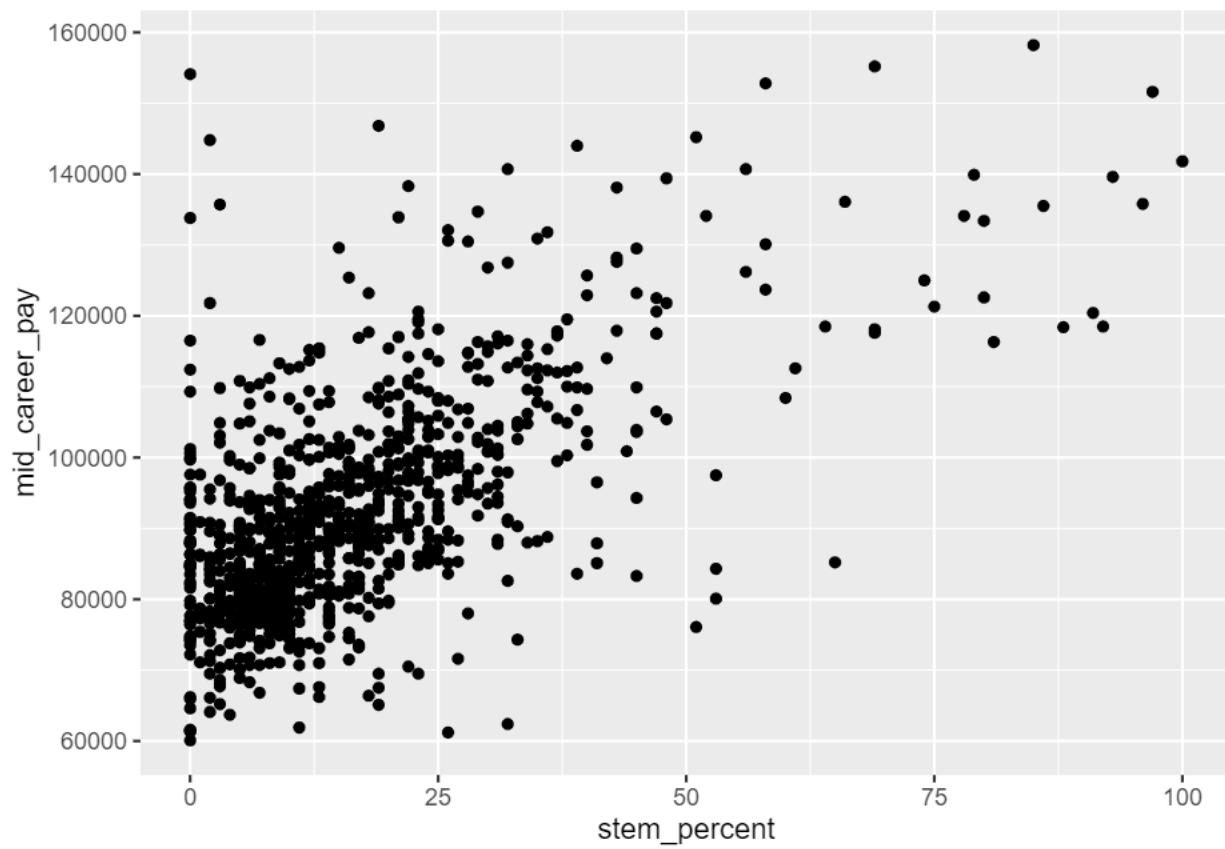
```
ggplot(sp, aes(stem_percent, early_career_pay)) + geom_point()
```



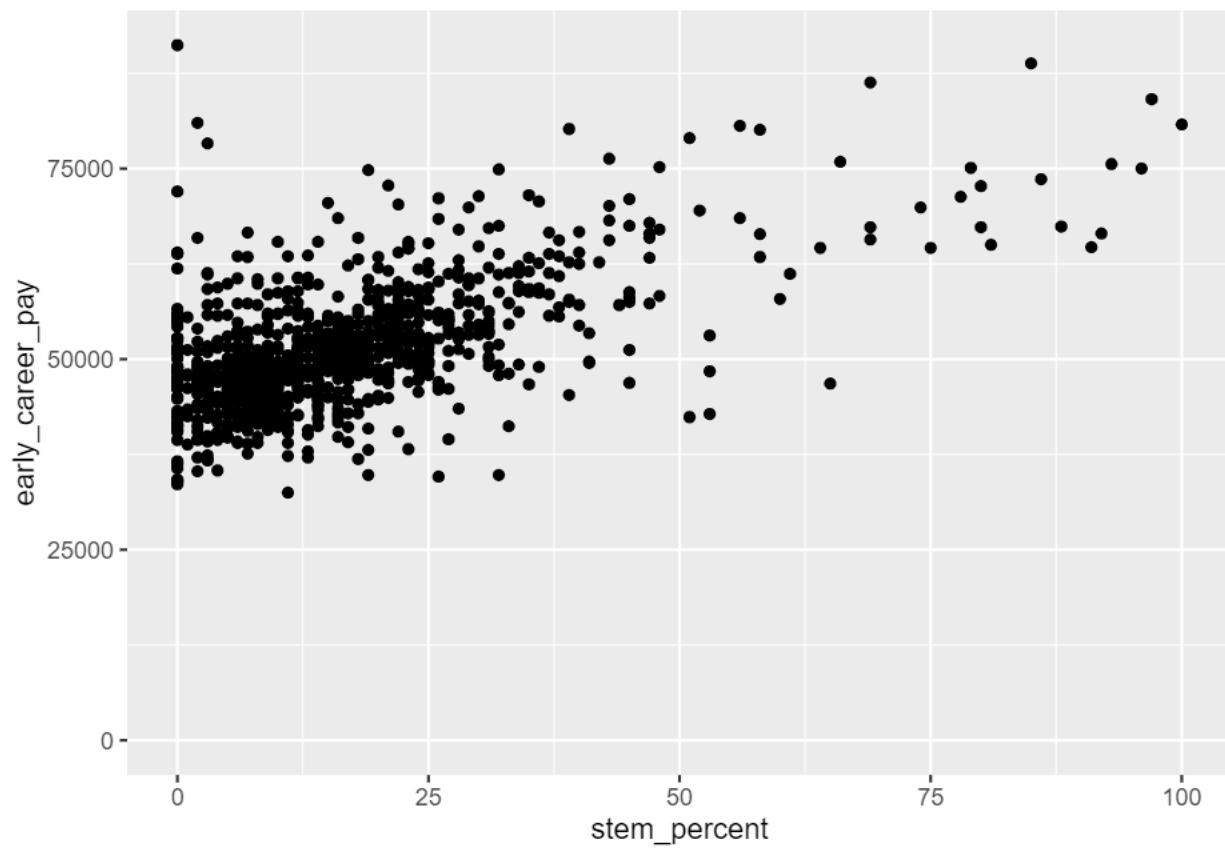
My conclusion is that there seems to be a weak positive correlation between these two variables.

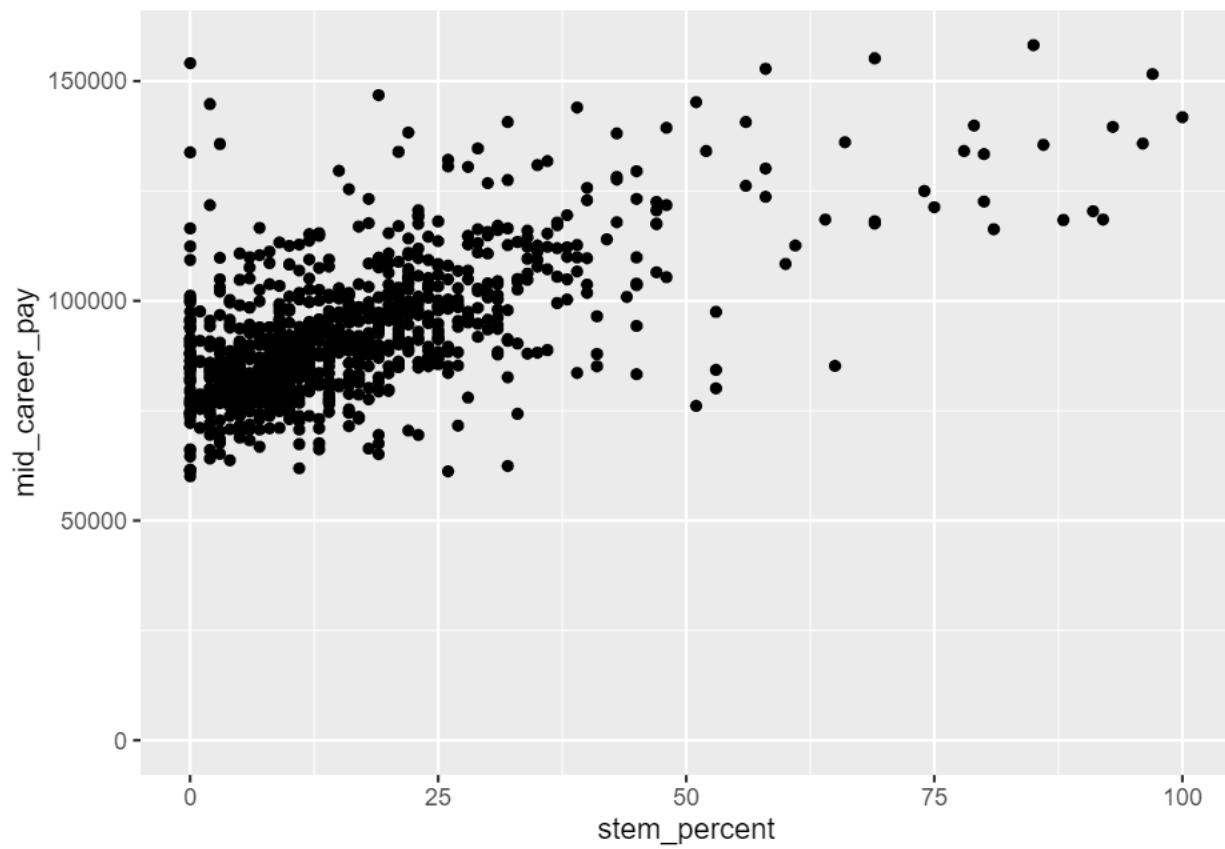
#In this graphic, perhaps the trend keeps?

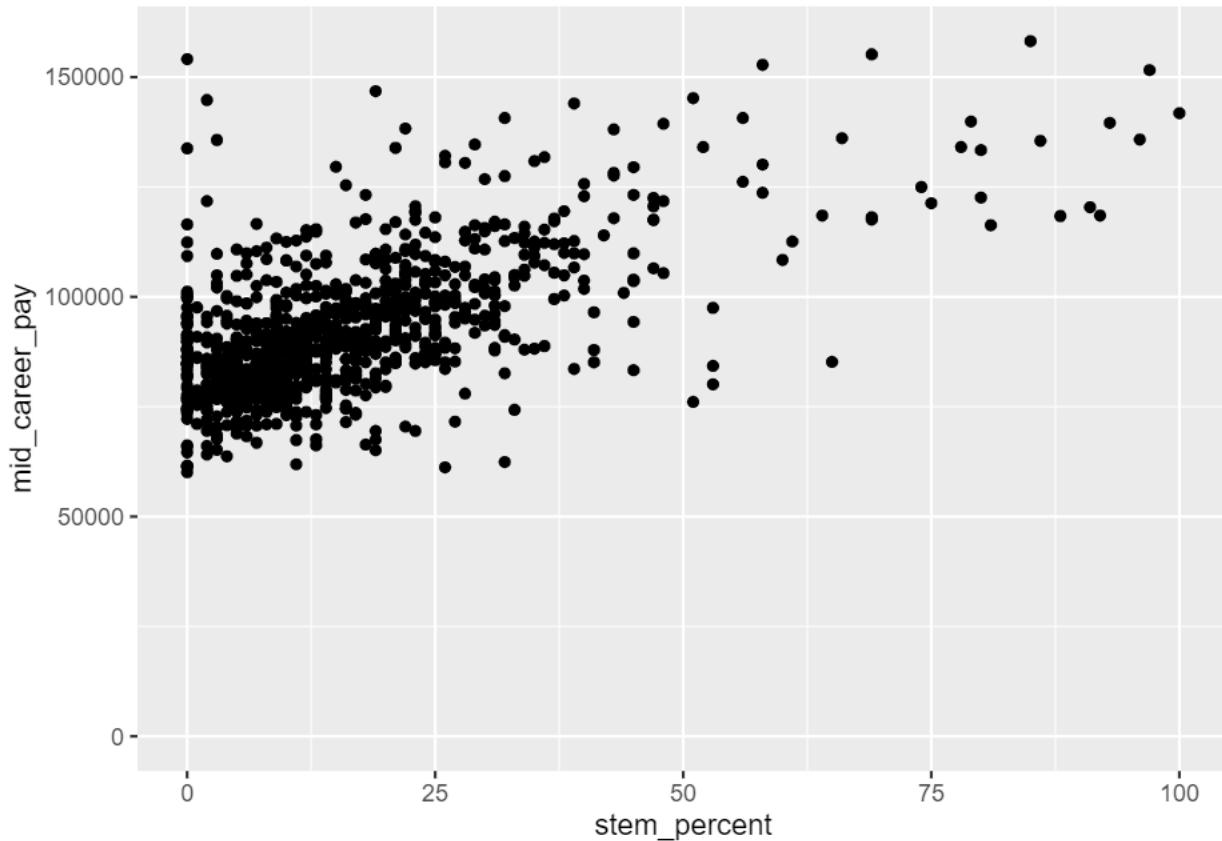
```
ggplot(sp, aes(stem_percent,mid_career_pay)) + geom_point()
```



#attempting to make the graphics easier to differentiate. I'm not sure what the difference is.
ggplot(sp, aes(x=stem_percent, y=early_career_pay)) + geom_point() + expand_limits(x=0, y=0)







```

income = ti %>%
  group_by(income_lvl) %>%
  summarize(Count=n()) %>%
  mutate(Percent = round((Count/sum(Count)*100))) %>%
  arrange(desc(Count))
income

## # A tibble: 5 x 3
##   income_lvl      Count Percent
##   <chr>          <int>   <dbl>
## 1 0 to 30,000    44969    22
## 2 30,001 to 48,000 43384    21
## 3 48_001 to 75,000 42600    20
## 4 75,001 to 110,000 40403    19
## 5 Over 110,000    37656    18

incomeByState = ti %>%
  group_by(income_lvl,state) %>%
  summarize(Count=n()) %>%
  mutate(Percent = round((Count/sum(Count)*100))) %>%
  arrange(desc(Count))

## `summarise()` has grouped output by 'income_lvl'. You can override using the
## `.`groups` argument.

incomeByState

## # A tibble: 255 x 4
## # Groups:   income_lvl [5]

```

```

##      income_lvl      state Count Percent
##      <chr>           <chr> <int>   <dbl>
## 1 0 to 30,000    NY     3460     8
## 2 30,001 to 48,000 NY     3357     8
## 3 48_001 to 75,000 NY     3333     8
## 4 0 to 30,000    CA     3290     7
## 5 75,001 to 110,000 NY     3125     8
## 6 30,001 to 48,000 CA     2909     7
## 7 Over 110,000   NY     2877     8
## 8 48_001 to 75,000 CA     2841     7
## 9 0 to 30,000    PA     2776     6
## 10 30,001 to 48,000 PA    2736     6
## # ... with 245 more rows

tcFacJoinSp = tcFactored %>%
  filter(tcFactored$type != 'For Profit') %>%
  inner_join(sp, by=c("name"="name"))

tcFacJoinSp


## # A tibble: 727 x 17
##      name      state state~1 type  degree~2 room_~3 in_st~4 in_st~5 out_o~6 out_o~7
##      <chr>     <chr> <chr> <chr> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 Adams St~ Colo~ CO    Publ~ 4 Year    8782    9440   18222   20456   29238
## 2 Adventis~ Flor~ FL    Priv~ 4 Year    4200    15150   19350   15150   19350
## 3 Agnes Sc~ Geor~ GA    Priv~ 4 Year   12330    41160   53490   41160   53490
## 4 Alabama ~ Alab~ AL    Publ~ 4 Year    5422    11068   16490   19396   24818
## 5 Alaska P~ Alas~ AK    Priv~ 4 Year    7300    20830   28130   20830   28130
## 6 Albany C~ New ~ NY    Priv~ 4 Year   10920    35105   46025   35105   46025
## 7 Albertus~ Conn~ CT    Priv~ 4 Year   13200    32060   45260   32060   45260
## 8 Albion C~ Mich~ MI    Priv~ 4 Year   12380    45775   58155   45775   58155
## 9 Alcorn S~ Miss~ MS    Publ~ 4 Year    9608    7144    16752    7144   16752
## 10 Allen Co~ Iowa IA   Priv~ 4 Year   7282    19970   27252   19970   27252
## # ... with 717 more rows, 7 more variables: degFactor <fct>, rank <dbl>,
## #   state_name <chr>, early_career_pay <dbl>, mid_career_pay <dbl>,
## #   make_world_better_percent <dbl>, stem_percent <dbl>, and abbreviated
## #   variable names 1: state_code, 2: degree_length, 3: room_and_board,
## #   4: in_state_tuition, 5: in_state_total, 6: out_of_state_tuition,
## #   7: out_of_state_total

ggplot(sp, aes(stem_percent,early_career_pay)) + geom_point() +
  geom_smooth(method="lm", se=FALSE) +
  ggtitle("Early career pay against prevalence of STEM majors in a school") +
  xlab("Prevalence of STEM majors (%)") +
  ylab("Early-career pay ($)") +
  theme_bw()

## `geom_smooth()` using formula 'y ~ x'

```

Early career pay against prevalence of STEM majors in a school



```
cor(sp$stem_percent, sp$early_career_pay)
```

```
## [1] 0.6050609
```

There seems to be a moderately strong correlation between the prevalence of STEM majors and how much graduates tend to

```
ECSPmodel = lm(early_career_pay ~ stem_percent, data=tcFacJoinSp)
summary(ECSPmodel)
```

```
##
## Call:
## lm(formula = early_career_pay ~ stem_percent, data = tcFacJoinSp)
##
## Residuals:
##     Min      1Q Median      3Q     Max 
## -21134   -4041   -624   3056  34761 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 45592.77    353.98 128.80 <2e-16 ***
## stem_percent  323.16     15.46  20.91 <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 6536 on 725 degrees of freedom
## Multiple R-squared:  0.3762, Adjusted R-squared:  0.3753 
## F-statistic: 437.2 on 1 and 725 DF,  p-value: < 2.2e-16
```

```
Slope: 323.30 Y-intercept: 45584.50
```

```
cor(x=sp$stem_percent,y=sp$early_career_pay)
```

```
## [1] 0.6050609
```

```
cor(x=sp$stem_percent,y=sp$mid_career_pay)
```

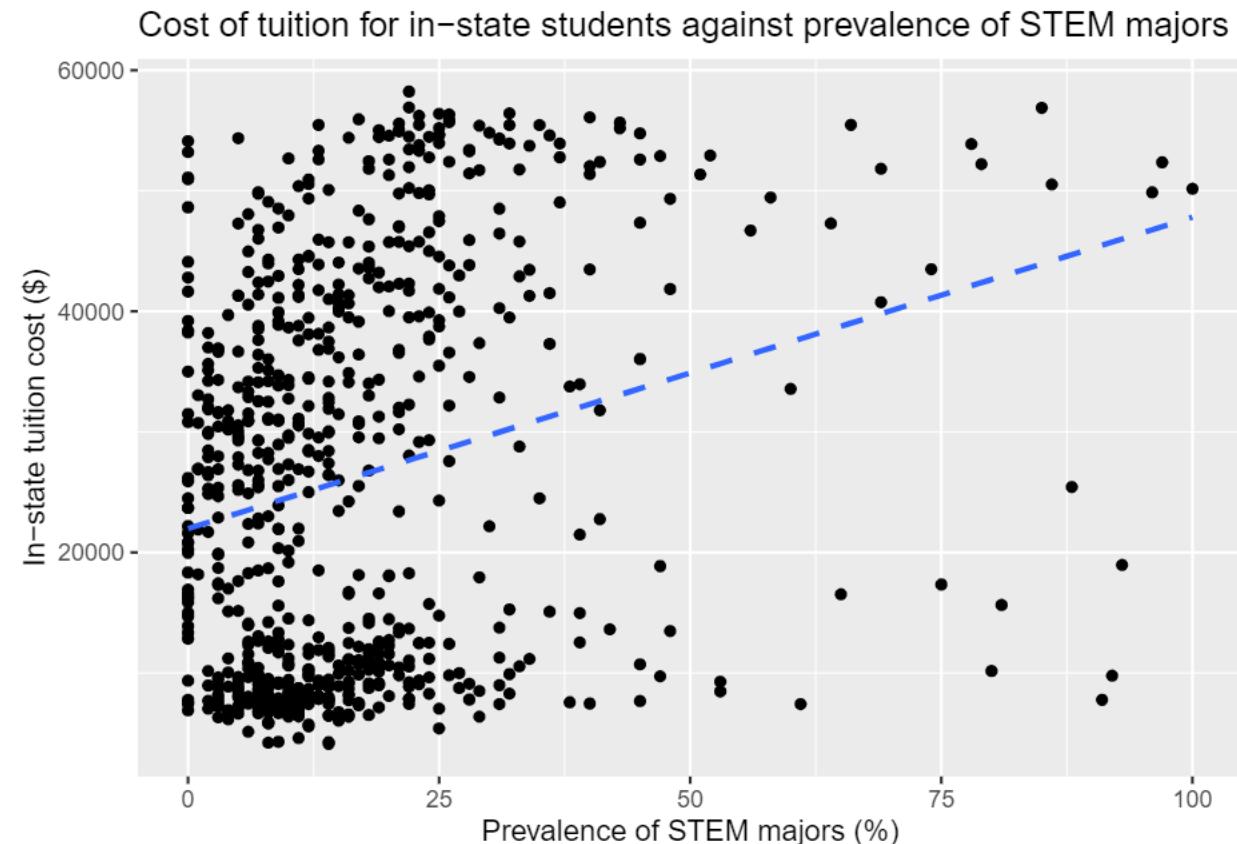
```
## [1] 0.6212143
```

Moderately strong positive association between prevalence of STEM majors at a school and general alumni early-career earnings.

Boring. STEM jobs tend to pay really well. Let's do something fun. Hear me out...

```
ggplot(tcFacJoinSp, aes(x=stem_percent,y=in_state_tuition)) +  
  geom_point() +  
  geom_smooth(method="lm",se=FALSE,lty=2) +  
  ggtitle("Cost of tuition for in-state students against prevalence of STEM majors") +  
  xlab("Prevalence of STEM majors (%)") +  
  ylab("In-state tuition cost ($)")
```

```
## `geom_smooth()` using formula 'y ~ x'
```

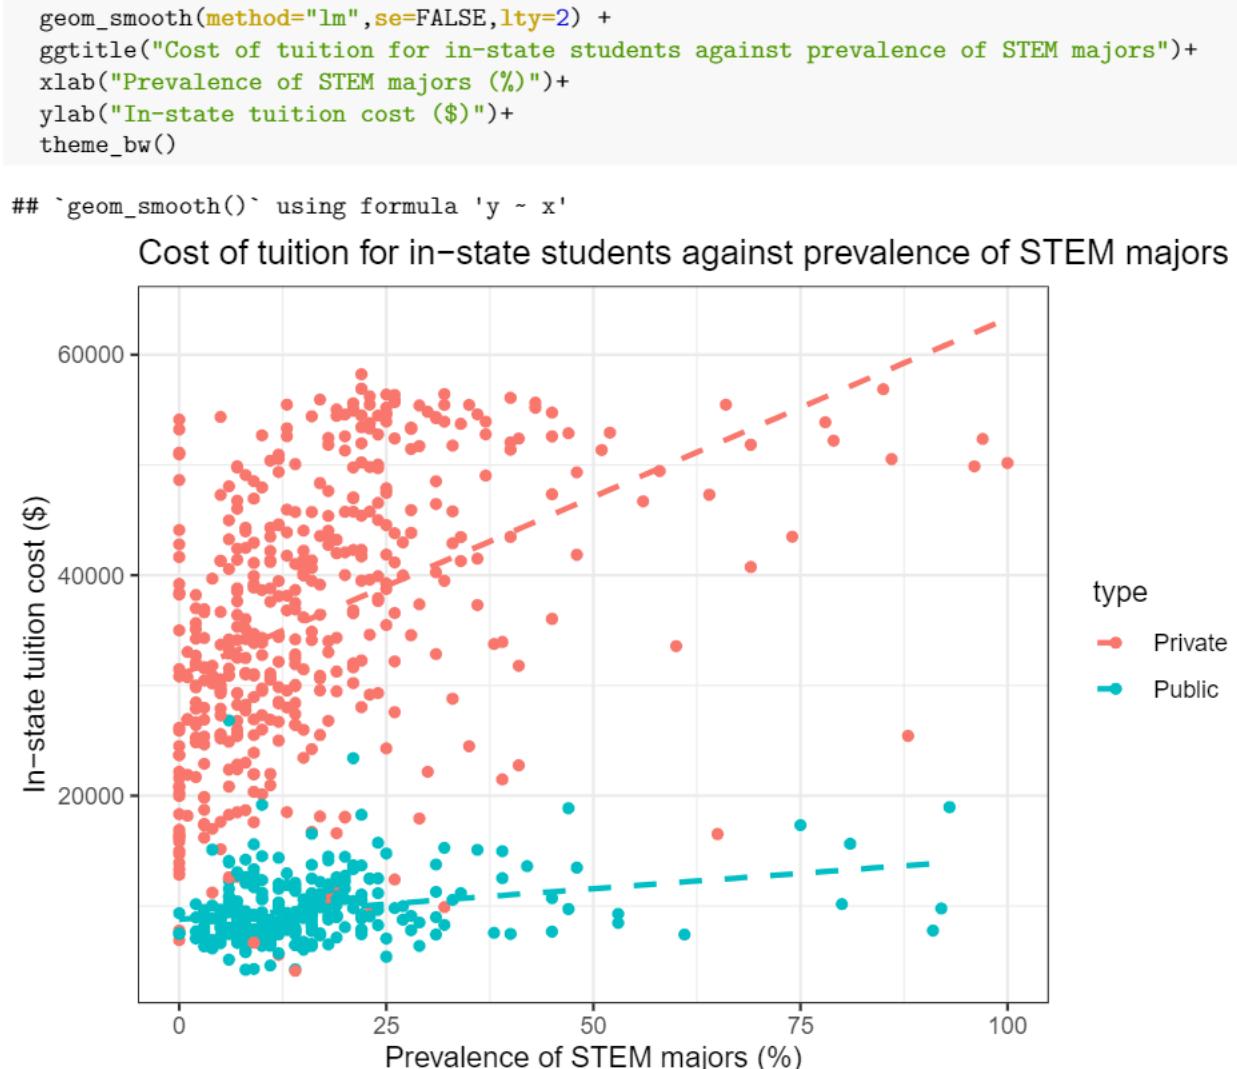


```
cor(x=tcFacJoinSp$stem_percent,y=tcFacJoinSp$in_state_tuition)
```

```
## [1] 0.2487816
```

Terrible fit. Couldn't predict the broad side of a barn!

```
ggplot(tcFacJoinSp, aes(x=stem_percent,y=in_state_tuition,color=type)) +  
  geom_point()
```



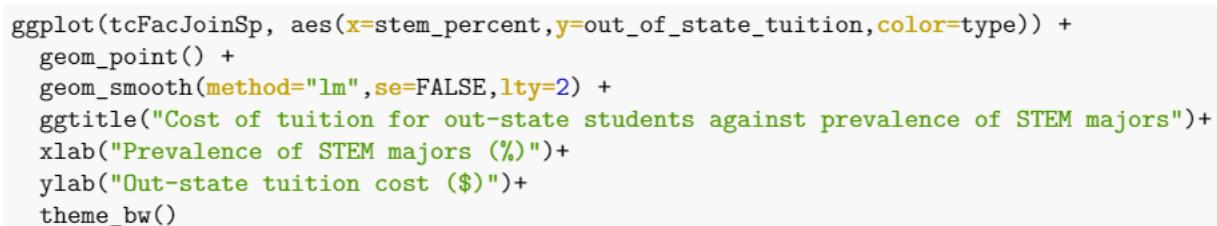
```

correlationsIS <- tcFacJoinSp %>%
  group_by(type) %>%
  summarise(r = cor(stem_percent, in_state_tuition))
correlationsIS

## # A tibble: 2 x 2
##   type     r
##   <chr>  <dbl>
## 1 Private 0.433
## 2 Public  0.265

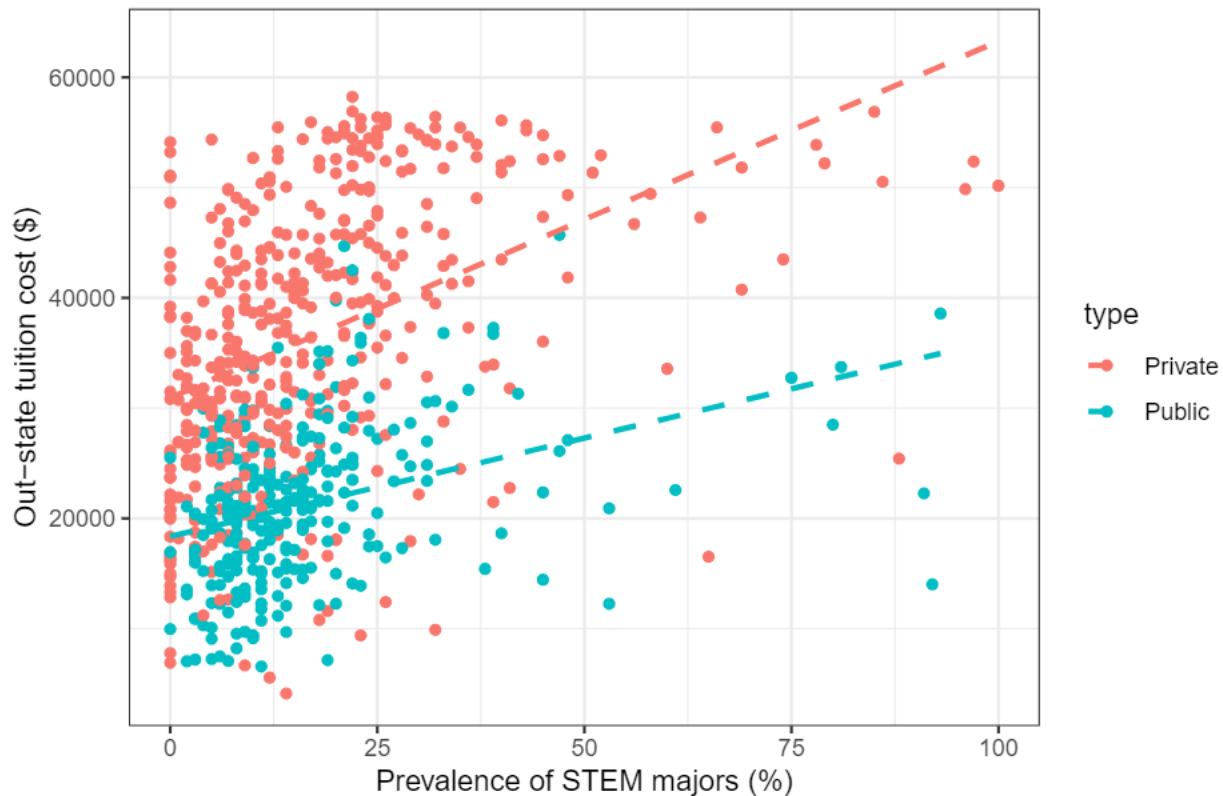
```

Moderately weak/moderate correlation between variables given institution type.



```
## `geom_smooth()` using formula 'y ~ x'
```

Cost of tuition for out-state students against prevalence of STEM majors



```
correlationsOOS <- tcFacJoinSp %>%
  group_by(type) %>%
  summarise(r = cor(stem_percent, out_of_state_tuition))
correlationsOOS
```

```
## # A tibble: 2 x 2
##   type     r
##   <chr>   <dbl>
## 1 Private 0.433
## 2 Public  0.356
```

Correlations show that there is a moderate correlation between the two variables given type.

These two graphs support the idea that tuition tends to be more expensive for schools that see greater proportions of STEM enrollment. Possibly due to the expensive equipment that generally comes with academic amenities like labs and associated equipment.

```
ISSPmodel=lm(in_state_tuition~stem_percent,data=tcFacJoinSp)
OSSPmodel=lm(out_of_state_tuition~stem_percent,data=tcFacJoinSp)
```

```
summary(ISSPmodel)
```

```
##
## Call:
## lm(formula = in_state_tuition ~ stem_percent, data = tcFacJoinSp)
##
## Residuals:
```

```

##      Min     1Q Median     3Q    Max
## -37709 -15280    1204  13279  32147
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 21970.98     855.47  25.683 < 2e-16 ***
## stem_percent   258.33      37.35   6.916 1.02e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 15790 on 725 degrees of freedom
## Multiple R-squared:  0.06189, Adjusted R-squared:  0.0606
## F-statistic: 47.83 on 1 and 725 DF, p-value: 1.02e-11
summary(OSSPmodel)

##
## Call:
## lm(formula = out_of_state_tuition ~ stem_percent, data = tcFacJoinSp)
##
## Residuals:
##      Min     1Q Median     3Q    Max
## -38392 -8930  -1347   9181  28263
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 25855.11     657.14   39.34 <2e-16 ***
## stem_percent   288.67      28.69   10.06 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12130 on 725 degrees of freedom
## Multiple R-squared:  0.1225, Adjusted R-squared:  0.1213
## F-statistic: 101.2 on 1 and 725 DF, p-value: < 2.2e-16

```